

# Logistic Regression on Amazon fine food dataset

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews> (<https://www.kaggle.com/snap/amazon-fine-food-reviews>)

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454

Number of users: 256,059

Number of products: 74,258

Timespan: Oct 1999 - Oct 2012

Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

**Objective:**

To perform Logistic regression using L2 regularization on different vectors like BOW, Tf-idf, Avg-W2vec & Tf-idf\_W2vec.

```
1 %matplotlib inline
2 import warnings
3 warnings.filterwarnings("ignore")
4
5 import sqlite3
6 import pandas as pd
7 import numpy as np
8 import nltk
9 import string
10 import matplotlib.pyplot as plt
11 import seaborn as sns
12 from sklearn.feature_extraction.text import TfidfTransformer
13 from sklearn.feature_extraction.text import TfidfVectorizer
14
15 from sklearn.feature_extraction.text import CountVectorizer
16 from sklearn.metrics import confusion_matrix
17 from sklearn import metrics
18 from sklearn.metrics import roc_curve, auc
19 from nltk.stem.porter import PorterStemmer
20
21 import re
22
23 import string
24 from nltk.corpus import stopwords
25 from nltk.stem import PorterStemmer
26 from nltk.stem.wordnet import WordNetLemmatizer
27
28 from gensim.models import Word2Vec
29 from gensim.models import KeyedVectors
30 import pickle
31
32 from tqdm import tqdm
33 import os
```

```
1 #Importing Train and test dataset
2 train_data=pd.read_csv("E:/Applied AI assignments/Amazon_fine_train_data.csv")
3 test_data=pd.read_csv("E:/Applied AI assignments/Amazon_fine_test_data.csv")
```

```
1 train_data=train_data.astype(str)
2 test_data=test_data.astype(str)
```

```
1 train_data.shape

(80000, 13)
```

```
1 train_data['Score'].value_counts()

positive    70407
negative    9593
Name: Score, dtype: int64
```

```
1 test_data.shape

(20000, 13)
```

```
1 test_data['Score'].value_counts()

positive    17322
negative     2678
Name: Score, dtype: int64
```

```
1 #Train data
2 y_train = train_data['Score']
3 x_train = train_data['CleanedText']
4
5 #Test data
6 y_test = test_data['Score']
7 x_test = test_data['CleanedText']
```

```
1 #Replacing Positive score with 0 and negative score with 1
2 y_train.replace('negative',1,inplace=True)
3 y_train.replace('positive',0,inplace=True)
4
5 y_test.replace('negative',1,inplace=True)
6 y_test.replace('positive',0,inplace=True)
```

```
1 from sklearn.linear_model import LogisticRegression
2 from sklearn.model_selection import RandomizedSearchCV
3 from sklearn.model_selection import TimeSeriesSplit
4 from sklearn.metrics import accuracy_score
5 from sklearn.metrics import recall_score
6 from sklearn.metrics import precision_score
7 from sklearn.metrics import f1_score
8 from sklearn.metrics import make_scorer
9 from sklearn.metrics import confusion_matrix
10 from sklearn.cross_validation import cross_val_score
11 from collections import Counter
12 from sklearn import cross_validation
13 from wordcloud import WordCloud
14 import matplotlib.pyplot as plt
15 from tqdm import tqdm
```

## Applying L2 regularization

### Randomised CV using L2

```

1 gamma_range = [0.000000001,0.000000001,0.00000001,0.0000001,0.000001,0.00001,0.0001,0.001,0.01,0.1,0.2,
2 T= TimeSeriesSplit(n_splits=5)
3 weight=[None,'balanced']
4
5 param_distributions = dict(C=gamma_range,class_weight=weight)
6 print(param_distributions)
7
8 # instantiate and fit the grid
9 grid = RandomizedSearchCV(LogisticRegression(penalty='l2'), param_distributions, cv=T, scoring='f1', re

```

```

{'C': [1e-10, 1e-09, 1e-08, 1e-07, 1e-06, 1e-05, 0.0001, 0.001, 0.01, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0, 5,
10, 15, 20, 30, 40, 50], 'class_weight': [None, 'balanced']}

```

## Binary Bow

```

1 count_vect = CountVectorizer(binary=True)
2
3 #Train data
4 vocabulary = count_vect.fit(x_train) #in scikit-learn
5 Bow_x_train= count_vect.transform(x_train)
6 print("the type of count vectorizer ",type(Bow_x_train))
7 print("the shape of out text BOW vectorizer ",Bow_x_train.get_shape())
8 print("the number of unique words ", Bow_x_train.get_shape()[1])

```

```

the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (80000, 33433)
the number of unique words 33433

```

```
1 #Test data
2 Bow_x_test = count_vect.transform(x_test)
3 print("the type of count vectorizer ",type(Bow_x_test))
4 print("the shape of out text BOW vectorizer ",Bow_x_test.get_shape())
5 print("the number of unique words ", Bow_x_test.get_shape()[1])
```

```
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (20000, 33433)
the number of unique words 33433
```

```
1 #Standardizing Bow_x_train and Bow_x_test
2 from sklearn.preprocessing import StandardScaler
3 Scaler=StandardScaler(with_mean=False)
4 Bow_x_train = Scaler.fit_transform(Bow_x_train)
5 Bow_x_test = Scaler.transform(Bow_x_test)
6
7 print(Bow_x_train.shape)
8 print(Bow_x_test.shape)
```

```
(80000, 33433)
(20000, 33433)
```


## Fitting Grid Search CV on BOW

```
1 grid.fit(Bow_x_train, y_train)
2
3 # examine the best model
4 print(grid.best_score_)
5 print(grid.best_params_)
```

```
0.6401387857438695
{'class_weight': 'balanced', 'C': 0.0001}
```

```
1 #Plotting C v/s CV_error
2 a=pd.DataFrame(grid.cv_results_)[['mean_test_score', 'std_test_score', 'params']]
3 a['C'] = [d.get('C') for d in a['params']]
4 b=a.sort_values(['C'])
5 CV_Error=1-b['mean_test_score']
6 C =b['C']
7
8
9 plt.plot(C,CV_Error)
10 plt.xlabel('C')
11 plt.ylabel('Cross-Validated Error')
```

```
Text(0,0.5,'Cross-Validated Error')
```



```
1
2  #{'class_weight': 'balanced', 'C': 0.0001}
3  LR_optimal=LogisticRegression(penalty='l2',C=0.0001,class_weight='balanced')
4
5  # fitting the model
6  LR_optimal.fit(Bow_x_train, y_train)
7
8  # predict the response
9  pred_bow = LR_optimal.predict(Bow_x_test)
10
11 # evaluate f1_score
12 f1_score = f1_score(y_test, pred_bow)
13
14 # Train & Test Error
15 print("The overall f1_score for the Train Data is : ", metrics.f1_score(y_train,LR_optimal.predict(Bow_
16 print("The overall f1_score for the Test Data is : ", metrics.f1_score(y_test,pred_bow))
```

The overall f1\_score for the Train Data is : 0.8089229382604777

The overall f1\_score for the Test Data is : 0.675998091299507

### Pertubation test

```
1  # Re-training the model after adding noise
2  Epsilon = np.random.normal(loc=0,scale =0.01)
3  Noise_Bow_x_train=Bow_x_train
4  Noise_Bow_x_train.data+=Epsilon
```

```
1  Noise_Bow_x_train.shape

(80000, 33433)
```



```
1
2 #{'class_weight': 'balanced', 'C': 0.0001}
3 LR_optimal_noise=LogisticRegression(penalty='l2',C=0.0001,class_weight='balanced')
4
5 # fitting the model
6 LR_optimal_noise.fit(Noise_Bow_x_train, y_train)
7
8 # predict the response
9 pred_bow = LR_optimal_noise.predict(Bow_x_test)
10
11 # evaluate f1_score
12 f1_score = f1_score(y_test, pred_bow)
13
14 # Train & Test Error
15 print("The overall f1_score for the Train Data is : ", metrics.f1_score(y_train,LR_optimal.predict(Noise_Bow_x_train)))
16 print("The overall f1_score for the Test Data is : ", metrics.f1_score(y_test,pred_bow))
```

The overall f1\_score for the Train Data is : 0.8088692595430168

The overall f1\_score for the Test Data is : 0.675998091299507

```
1 #Features
2 feature_names = np.array(vocabulary.get_feature_names())
3 feature_names.shape
```

(33433,)

```
1 #Weights before adding noise
2 LR_optimal.coef_.shape
```

(1, 33433)

```
1 #Weights after adding noise  
2 LR_optimal_noise.coef_.shape  
  
(1, 33433)
```

```

1 merge_arr = np.concatenate([LR_optimal.coef_, LR_optimal_noise.coef_], axis=0)
2 merge=pd.DataFrame(data=merge_arr,columns=feature_names).transpose()
3 merge[2]=((merge[1]-merge[0])/merge[0])*100
4 merge

```

	0	1	2
aaa	-0.001722	-0.001722	0.002999
aaaaaaaaagghh	-0.002087	-0.002087	-0.012497
aaaaah	-0.002006	-0.002006	0.005672
aaaaahhhhhhhhhhhhhhh	-0.000765	-0.000765	0.000203
aaaah	-0.000964	-0.000964	-0.003298
aaah	-0.001224	-0.001224	-0.004986
aachen	0.005927	0.005927	-0.008612
aad	-0.000364	-0.000364	-0.071263
aadp	-0.001017	-0.001017	-0.024636
aafco	-0.000879	-0.000878	-0.090446
aagh	-0.003549	-0.003549	-0.004424
aah	-0.002004	-0.002004	0.000561
aahh	-0.001183	-0.001182	-0.024529
aand	-0.001929	-0.001929	-0.005073
aardvark	-0.003908	-0.003909	0.014199
aarrgh	0.009679	0.009679	0.000378
ab	-0.002795	-0.002795	0.003201
aback	-0.006500	-0.006501	0.003168
abandon	0.003590	0.003590	-0.012733
abaolut	-0.001443	-0.001442	-0.028476
abattoir	-0.000879	-0.000880	0.026223
abba	-0.002686	-0.002686	-0.004696
abbey	-0.002013	-0.002013	0.001276
abbi	-0.003064	-0.003064	-0.006595

	0	1	2
abbott	-0.000479	-0.000479	-0.044658
abbrevi	-0.000541	-0.000541	-0.002948
abc	0.000694	0.000695	0.252658
abcstor	-0.001687	-0.001687	0.020410
abd	-0.001774	-0.001774	-0.028242
abdomen	0.000109	0.000109	-0.001692
...	...	...	...
zot	-0.008788	-0.008788	0.005860
zotz	-0.004918	-0.004919	0.003821
zour	0.002518	0.002518	0.000398
zout	-0.001946	-0.001946	-0.022958
zowi	-0.000718	-0.000718	-0.045905
zreport	-0.004741	-0.004741	0.000733
zsweet	-0.002095	-0.002094	-0.028194
zuc	-0.002182	-0.002182	-0.009051
zucchini	0.004017	0.004017	0.002359
zuccini	-0.004626	-0.004625	-0.005549
zuccnini	-0.000321	-0.000321	-0.013392
zuchinni	-0.003164	-0.003164	0.001102
zuke	-0.001739	-0.001739	-0.015763
zulu	-0.001037	-0.001037	-0.007865
zum	-0.000321	-0.000320	-0.059621
zummi	-0.000321	-0.000320	-0.059621
zune	-0.003632	-0.003632	-0.002227
zupreem	-0.000610	-0.000610	-0.001370
zurich	-0.001767	-0.001767	-0.010648
zwar	-0.000111	-0.000111	0.007019
zwieback	0.003339	0.003340	0.018697
zwiebeck	-0.002276	-0.002275	-0.025352

	0	1	2
<b>zydeco</b>	-0.001672	-0.001672	0.003113
<b>zzzzzs</b>	-0.002309	-0.002309	0.006540
<b>zzzzzz</b>	-0.000076	-0.000076	-0.063760
<b>zzzzzzzz</b>	0.007196	0.007195	-0.003373
<b>zzzzzzzzzz</b>	-0.000630	-0.000630	0.014911
<b>zzzzzzzzzzzz</b>	-0.000379	-0.000379	-0.081017
<b>zzzzzzzzzzzzzz</b>	-0.002822	-0.002822	0.004119
<b>çay</b>	-0.001229	-0.001229	-0.016318

33433 rows x 3 columns

```
1 merge[merge[2]>30].shape
```

(3, 3)

3 features out of 33433 shows percentage change > 30 post pertubation test i.e 0.0089%

We can say that our data isn't much affected by multicollinearity

```
1 feature_names = np.array(vocabulary.get_feature_names())
2 sorted_coef_index = LR_optimal.coef_[0].argsort()
```

```
1 #Top 20 positive features
2 p=feature_names[sorted_coef_index[:20]]
3
4 sp = ""
5 for i in p:
6     sp += str(i)+", "
7 print(sp)
```

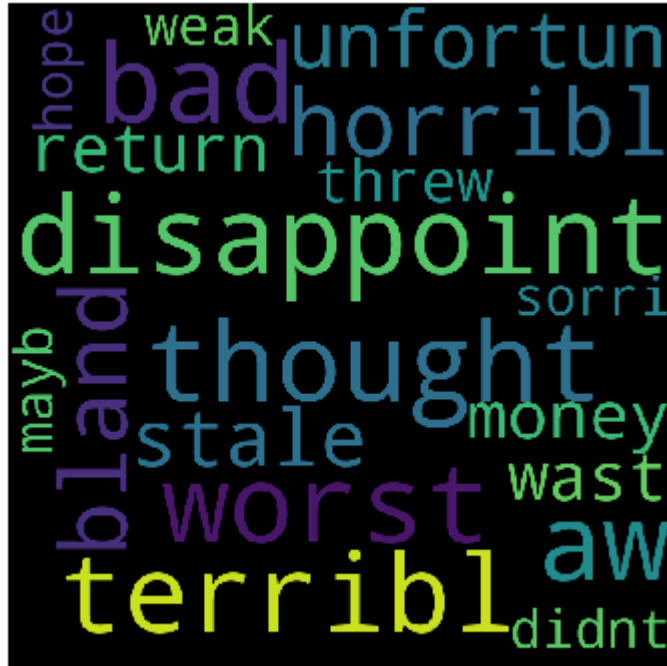
great, love, best, delici, perfect, excel, good, favorit, nice, wonder, find, tasti, easi, amaz, thank, addict, alway, keep, year, snack,

```
1 n=feature_names[sorted_coef_index[: -21: -1]]
2
3 sn = ""
4 for i in n:
5     sn += str(i) + ", "
6 print(sn)
```

disappoint, worst, terrible, thought, bad, aw, horrible, bland, unfortunate, stale, would, money, return, waste, threw, didn't, maybe, weak, sorry, hope,

```
1 print("***** Top 20 Negative words *****")
2 wordcloud = WordCloud(width = 800, height = 800,
3                       background_color = 'black',
4                       min_font_size = 10).generate(sn)
5
6 # plot the WordCloud image
7 plt.figure(figsize = (5,5), facecolor = None)
8 plt.imshow(wordcloud)
9 plt.axis("off")
10 plt.tight_layout(pad = 0)
11 plt.show()
12
13
14 print("***** Top 20 Positive words *****")
15 wordcloud = WordCloud(width = 800, height = 800,
16                       background_color = 'black',
17                       min_font_size = 10).generate(sp)
18
19 # plot the WordCloud image
20 plt.figure(figsize = (5,5), facecolor = None)
21 plt.imshow(wordcloud)
22 plt.axis("off")
23 plt.tight_layout(pad = 0)
24 plt.show()
```

\*\*\*\*\* Top 20 Negative words \*\*\*\*\*



\*\*\*\*\* Top 20 Positive words \*\*\*\*\*







```
1 #Confusion matrix
2 C = confusion_matrix(y_test, pred_bow)
3 A = (((C.T)/(C.sum(axis=1))).T)
4 B = (C/C.sum(axis=0))
5 labels = [0,1]
```

```
1 print("-"*20, "Confusion matrix", "-"*20)
2 plt.figure(figsize=(12,7))
3 sns.heatmap(C, annot=True, cmap="YlGnBu", fmt="g", xticklabels=labels, yticklabels=labels)
4 plt.xlabel('Predicted Class')
5 plt.ylabel('Original Class')
6 plt.show()
7
8 print("-"*20, "Precision matrix (Column Sum=1)", "-"*20)
9 plt.figure(figsize=(12,7))
10 sns.heatmap(B, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels, yticklabels=labels)
11 plt.xlabel('Predicted Class')
12 plt.ylabel('Original Class')
13 plt.show()
14
15     # representing B in heatmap format
16 print("-"*20, "Recall matrix (Row sum=1)", "-"*20)
17 plt.figure(figsize=(12,7))
18 sns.heatmap(A, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels, yticklabels=labels)
19 plt.xlabel('Predicted Class')
20 plt.ylabel('Original Class')
21 plt.show()
```

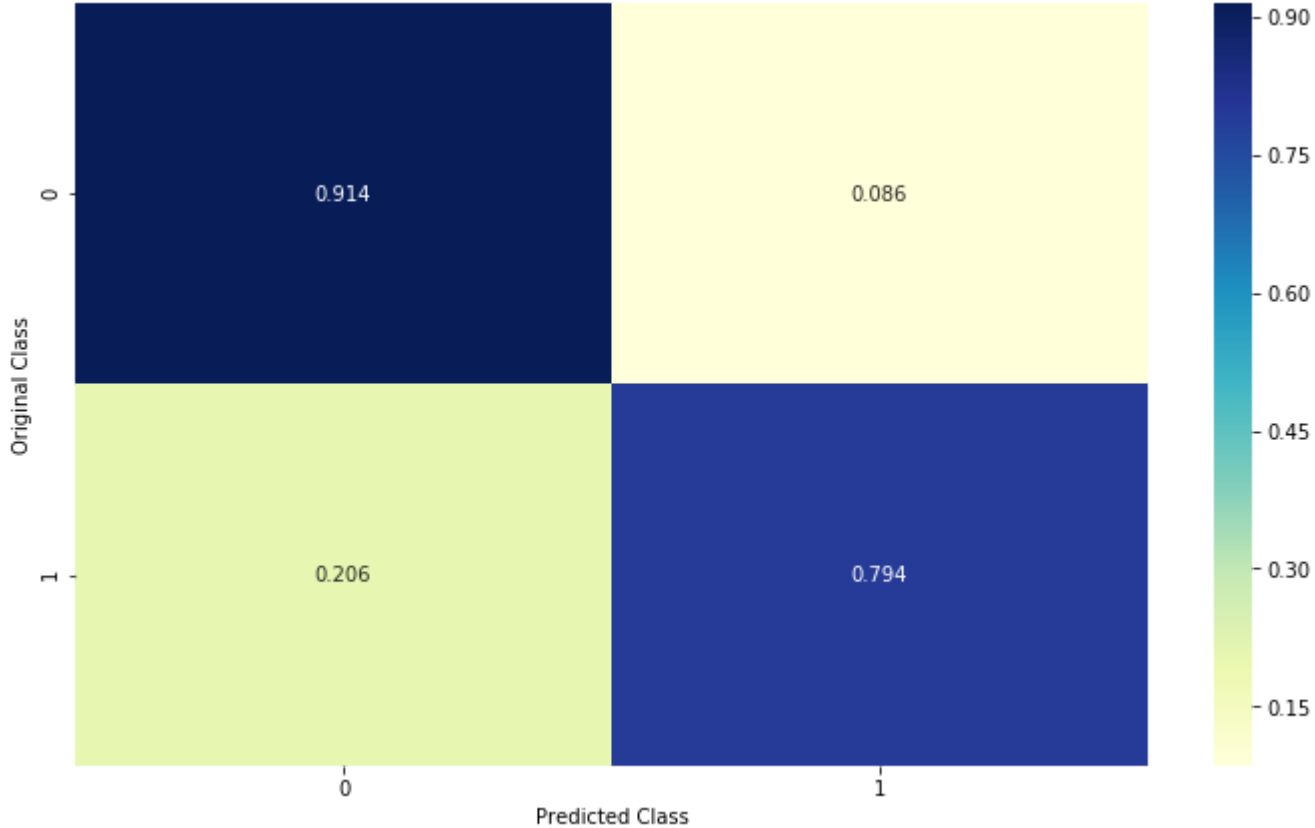
----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



Tf-idf

```
1 #Initiating Vectorizer
2 count_vect = TfidfVectorizer(ngram_range=(1,2))
3
4 #Train data
5 vocabulary = count_vect.fit(x_train)
6 Tfidf_x_train= count_vect.transform(x_train)
7 print("the type of count vectorizer ",type(Tfidf_x_train))
8 print("the shape of out text BOW vectorizer ",Tfidf_x_train.get_shape())
9 print("the number of unique words ", Tfidf_x_train.get_shape()[1])
```

the type of count vectorizer <class 'scipy.sparse.csr.csr\_matrix'>  
the shape of out text BOW vectorizer (80000, 1013943)  
the number of unique words 1013943

```
1 #Test data
2 Tfidf_x_test= count_vect.transform(x_test)
3 print("the type of count vectorizer ",type(Tfidf_x_test))
4 print("the shape of out text BOW vectorizer ",Tfidf_x_test.get_shape())
5 print("the number of unique words ", Tfidf_x_test.get_shape()[1])
```

the type of count vectorizer <class 'scipy.sparse.csr.csr\_matrix'>  
the shape of out text BOW vectorizer (20000, 1013943)  
the number of unique words 1013943

```
1 #Standardizing
2 from sklearn.preprocessing import StandardScaler
3 Standard=StandardScaler(with_mean=False)
4 Tfidf_x_train = Standard.fit_transform(Tfidf_x_train)
5 Tfidf_x_test = Standard.transform(Tfidf_x_test)
6
7 print(Tfidf_x_train.shape)
8 print(Tfidf_x_test.shape)
```

(80000, 1013943)  
(20000, 1013943)

## Fitting Randomsearch on Tf-Idf


```
1 grid.fit(Tfidf_x_train, y_train)
2
3 # examine the best model
4 print(grid.best_score_)
5 print(grid.best_params_)
```

```
0.5390181835332085
```

```
{'class_weight': 'balanced', 'C': 1e-07}
```

```
1 #Plotting C v/s CV_error
2 a=pd.DataFrame(grid.cv_results_)[['mean_test_score', 'std_test_score', 'params']]
3 a['C'] = [d.get('C') for d in a['params']]
4 b=a.sort_values(['C'])
5 CV_Error=1-b['mean_test_score']
6 C =b['C']
7
8
9 plt.plot(C,CV_Error)
10 plt.xlabel('C')
11 plt.ylabel('Cross-Validated Error')
```

```
Text(0,0.5,'Cross-Validated Error')
```





```
1
2 #{'class_weight': 'balanced', 'C': 1e-07}
3 LR_optimal=LogisticRegression(penalty='l2',C=0.0000001,class_weight='balanced')
4
5 # fitting the model
6 LR_optimal.fit(Tfidf_x_train, y_train)
7
8 # predict the response
9 pred_tfidf = LR_optimal.predict(Tfidf_x_test)
10
11 # evaluate accuracy
12 f1_score = f1_score(y_test, pred_tfidf)
13
14 # Train & Test Error
15 print("The overall f1_score for the Train Data is : ", metrics.f1_score(y_train,LR_optimal.predict(Tfidf_x_train)))
16 print("The overall f1_score for the Test Data is : ", metrics.f1_score(y_test,pred_tfidf))
```

The overall f1\_score for the Train Data is : 0.9030708364732479

The overall f1\_score for the Test Data is : 0.5339728217426059

### Pertubation test

```
1 # Re-training the model after adding noise
2 Epsilon = np.random.normal(loc=0,scale =0.01)
3 Noise_Tfidf_x_train=Tfidf_x_train
4 Noise_Tfidf_x_train.data+=Epsilon
```

```
1
2 #{'class_weight': 'balanced', 'C': 1e-07}
3 LR_optimal_noise=LogisticRegression(penalty='l2',C=0.000001,class_weight='balanced')
4
5 # fitting the model
6 LR_optimal_noise.fit(Noise_Tfidf_x_train, y_train)
7
8 # predict the response
9 pred_tfidf = LR_optimal_noise.predict(Tfidf_x_test)
10
11 # evaluate accuracy
12 f1_score = f1_score(y_test, pred_tfidf)
13
14 # Train & Test Error
15 print("The overall f1_score for the Train Data is : ", metrics.f1_score(y_train,LR_optimal_noise.predict(y_train)))
16 print("The overall f1_score for the Test Data is : ", metrics.f1_score(y_test,pred_tfidf))
```

The overall f1\_score for the Train Data is : 0.9030283049969386

The overall f1\_score for the Test Data is : 0.5340948029697316

```
1 #Features
2 feature_names = np.array(vocabulary.get_feature_names())
3 feature_names.shape
```

(1013943,)

```
1 LR_optimal.coef_.shape
```

(1, 1013943)

```
1 LR_optimal_noise.coef_.shape
```

(1, 1013943)

```

1 merge_arr = np.concatenate([LR_optimal.coef_, LR_optimal_noise.coef_], axis=0)
2 merge=pd.DataFrame(data=merge_arr,columns=feature_names).transpose()
3 merge[2]=((merge[1]-merge[0])/merge[0])*100
4 merge

```

	0	1	2
aaa	-0.000010	-0.000010	0.003393
aaa condit	-0.000008	-0.000008	0.002404
aaa perfect	-0.000008	-0.000008	0.000782
aaaaaaaagghh	-0.000008	-0.000008	0.000995
aaaaah	-0.000011	-0.000011	0.002724
aaaaah awak	-0.000008	-0.000008	0.000999
aaaaah satisfi	-0.000008	-0.000008	0.002002
aaaaahhhhhhhhhhhhhhhhh	-0.000008	-0.000008	0.001810
aaaaahhhhhhhhhhhhhhhhh angel	-0.000008	-0.000008	0.001810
aaaah	-0.000007	-0.000007	-0.000594
aaaah snob	-0.000007	-0.000007	-0.000594
aaah	-0.000013	-0.000013	0.002666
aaah inhal	-0.000008	-0.000008	0.000134
aaah miss	-0.000008	-0.000008	0.000481
aaah sip	-0.000008	-0.000008	0.000763
aachen	0.000045	0.000045	0.000334
aachen munich	0.000045	0.000045	0.000334
aad	-0.000008	-0.000008	-0.001176
aad sausag	-0.000008	-0.000008	-0.001176
aadp	-0.000008	-0.000008	0.000288
aafco	0.000002	0.000002	0.041650
aafco also	-0.000007	-0.000007	-0.002049
aafco certifi	0.000018	0.000018	0.001439
aafco countri	-0.000007	-0.000007	-0.002049

	0	1	2
aafco definit	0.000042	0.000042	0.001049
aafco dog	-0.000008	-0.000008	0.000260
aafco guidelin	-0.000008	-0.000008	-0.000410
aafco requir	-0.000008	-0.000008	0.001271
aagh	-0.000008	-0.000008	0.002790
aagh yelp	-0.000008	-0.000008	0.002790
...	...	...	...
zum heal	-0.000008	-0.000008	0.000441
zummi	-0.000008	-0.000008	0.000441
zummi love	-0.000008	-0.000008	0.000441
zummi tast	-0.000008	-0.000008	0.000441
zummi tri	-0.000008	-0.000008	0.000441
zune	-0.000008	-0.000008	0.002882
zune video	-0.000008	-0.000008	0.002882
zupreem	-0.000008	-0.000008	0.000569
zupreem ferret	-0.000008	-0.000008	0.000569
zurich	-0.000008	-0.000008	0.001963
zurich schnatzlet	-0.000008	-0.000008	0.001963
zwar	-0.000007	-0.000007	0.002107
zwar billig	-0.000007	-0.000007	0.002107
zwieback	0.000022	0.000022	0.004092
zwieback toast	0.000022	0.000022	0.004092
zwiebeck	-0.000008	-0.000008	0.001889
zwiebeck toast	-0.000008	-0.000008	0.001889
zydeco	-0.000008	-0.000008	0.001634
zydeco saturday	-0.000008	-0.000008	0.001634
zzzzzs	-0.000011	-0.000011	0.002403
zzzzzs larg	-0.000008	-0.000008	0.002258
zzzzzz	-0.000007	-0.000007	-0.001191

	0	1	2
zzzzzz say	-0.000007	-0.000007	-0.001191
zzzzzzzz	0.000053	0.000053	0.001704
zzzzzzzz high	0.000053	0.000053	0.001704
zzzzzzzzzz	-0.000007	-0.000007	0.002698
zzzzzzzzzzzz	-0.000008	-0.000008	-0.000450
zzzzzzzzzzzz final	-0.000008	-0.000008	-0.000450
zzzzzzzzzzzzzz	-0.000008	-0.000008	0.001947
çay	-0.000008	-0.000008	0.000097

1013943 rows x 3 columns

```
1 merge[merge[2]>30].shape
```

(1, 3)

1 features out of 1013943 shows percentage change > 30 post pertubation test i.e 0%

We can say that our data isn't affected by multicollinearity

```
1 feature_names = np.array(vocabulary.get_feature_names())
2 sorted_coef_index = LR_optimal.coef_[0].argsort()
```

```
1 #Top 20 positive features
2 p=feature_names[sorted_coef_index[:20]]
3
4 sp = ""
5 for i in p:
6     sp += str(i)+", "
7 print(sp)
```


great, love, best, delici, perfect, favorit, good, find, make, high recommend, easi, excel, wonder, nice, use, snack, alway, keep, add, tasti,

```
1 n=feature_names[sorted_coef_index[:-21:-1]]  
2  
3 sn = ""  
4 for i in n:  
5     sn += str(i)+", "  
6 print(sn)
```

disappoint,wast money,worst,wast,aw,horribl,terribl,return,threw,money,bad,refund,stale,wont buy,thought,mayb,didnt,dont wa  
st,away,wors,

```
1 print("***** Top 20 Negative words *****")
2 wordcloud = WordCloud(width = 800, height = 800,
3                       background_color = 'black',
4                       min_font_size = 10).generate(sn)
5
6 # plot the WordCloud image
7 plt.figure(figsize = (5,5), facecolor = None)
8 plt.imshow(wordcloud)
9 plt.axis("off")
10 plt.tight_layout(pad = 0)
11 plt.show()
12
13
14 print("***** Top 20 Positive words *****")
15 wordcloud = WordCloud(width = 800, height = 800,
16                       background_color = 'black',
17                       min_font_size = 10).generate(sp)
18
19 # plot the WordCloud image
20 plt.figure(figsize = (5,5), facecolor = None)
21 plt.imshow(wordcloud)
22 plt.axis("off")
23 plt.tight_layout(pad = 0)
24 plt.show()
```

\*\*\*\*\* Top 20 Negative words \*\*\*\*\*



\*\*\*\*\* Top 20 Positive words \*\*\*\*\*







```
1 #Confusion matrix
2 C = confusion_matrix(y_test, pred_tfidf)
3 A = (((C.T)/(C.sum(axis=1))).T)
4 B = (C/C.sum(axis=0))
5 labels = [0,1]
```

```
1  print("-"*20, "Confusion matrix", "-"*20)
2  plt.figure(figsize=(12,7))
3  sns.heatmap(C, annot=True, cmap="YlGnBu", fmt="g", xticklabels=labels, yticklabels=labels)
4  plt.xlabel('Predicted Class')
5  plt.ylabel('Original Class')
6  plt.show()
7
8  print("-"*20, "Precision matrix (Column Sum=1)", "-"*20)
9  plt.figure(figsize=(12,7))
10 sns.heatmap(B, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels, yticklabels=labels)
11  plt.xlabel('Predicted Class')
12  plt.ylabel('Original Class')
13  plt.show()
14
15  # representing B in heatmap format
16  print("-"*20, "Recall matrix (Row sum=1)", "-"*20)
17  plt.figure(figsize=(12,7))
18  sns.heatmap(A, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels, yticklabels=labels)
19  plt.xlabel('Predicted Class')
20  plt.ylabel('Original Class')
21  plt.show()
```


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



### Avg-W2Vec

```
1 #W2V List of Training data
2 i=0
3 list_of_sent_train=[]
4 for sent in train_data['CleanedText'].values:
5     list_of_sent_train.append(sent.split())
```

```
1 #W2V List of Test data
2 i=0
3 list_of_sent_test=[]
4 for sent in test_data['CleanedText'].values:
5     list_of_sent_test.append(sent.split())
```

```
1 #Training W2V train model
2 # min_count = 5 considers only words that occurred at least 5 times
3 w2v_model_train=Word2Vec(list_of_sent_train,min_count=5,size=50, workers=6)
```

```
1 w2v_words_train = list(w2v_model_train.wv.vocab)
2 print("number of words that occurred minimum 5 times ",len(w2v_words_train))
3 print("sample words ", w2v_words_train[0:50])
```

number of words that occurred minimum 5 times 11361

sample words ['witti', 'littl', 'book', 'make', 'son', 'laugh', 'loud', 'car', 'drive', 'along', 'alway', 'sing', 'refrain', 'hes', 'learn', 'whale', 'india', 'droop', 'love', 'new', 'word', 'introduc', 'silli', 'classic', 'will', 'bet', 'still', 'abl', 'memori', 'colleg', 'rememb', 'see', 'show', 'air', 'televis', 'year', 'ago', 'child', 'sister', 'later', 'bought', 'day', 'thirti', 'someth', 'use', 'seri', 'song', 'student', 'teach', 'preschool']

```
1  #Train data
2  # average Word2Vec
3  # compute average word2vec for each review.
4  sent_vectors_train_avgw2v = []; # the avg-w2v for each sentence/review is stored in this list
5  for sent in list_of_sent_train: # for each review/sentence
6      sent_vec = np.zeros(50) # as word vectors are of zero length
7      cnt_words = 0; # num of words with a valid vector in the sentence/review
8      for word in sent: # for each word in a review/sentence
9          if word in w2v_words_train:
10             vec = w2v_model_train.wv[word]
11             sent_vec += vec
12             cnt_words += 1
13         if cnt_words != 0:
14             sent_vec /= cnt_words
15         sent_vectors_train_avgw2v.append(sent_vec)
16 print(len(sent_vectors_train_avgw2v))
17 print(len(sent_vectors_train_avgw2v[0]))
```

80000

50

```
1 #Test data
2 # average Word2Vec
3 # compute average word2vec for each review.
4 sent_vectors_test_avgw2v = []; # the avg-w2v for each sentence/review is stored in this list
5 for sent in list_of_sent_test: # for each review/sentence
6     sent_vec = np.zeros(50) # as word vectors are of zero length
7     cnt_words = 0; # num of words with a valid vector in the sentence/review
8     for word in sent: # for each word in a review/sentence
9         if word in w2v_words_train:
10             vec = w2v_model_train.wv[word]
11             sent_vec += vec
12             cnt_words += 1
13     if cnt_words != 0:
14         sent_vec /= cnt_words
15     sent_vectors_test_avgw2v.append(sent_vec)
16 print(len(sent_vectors_test_avgw2v))
17 print(len(sent_vectors_test_avgw2v[0]))
```

20000

50

```
1 #Standardizing Avg-W2v
2 from sklearn.preprocessing import StandardScaler
3
4 Standard=StandardScaler()
5 sent_vectors_train_avgw2v = Standard.fit_transform(sent_vectors_train_avgw2v)
6 sent_vectors_test_avgw2v = Standard.transform(sent_vectors_test_avgw2v)
7
8 print(sent_vectors_train_avgw2v.shape)
9 print(sent_vectors_test_avgw2v.shape)
```

(80000, 50)

(20000, 50)

### Fitting grid search on Avg-W2V




```
1 grid.fit(sent_vectors_train_avgw2v, y_train)
2
3 # examine the best model
4 print(grid.best_score_)
5 print(grid.best_params_)
```

0.5134126417862019

{'class\_weight': 'balanced', 'C': 50}

```
1 #Plotting C v/s CV_error
2 a=pd.DataFrame(grid.cv_results_)[['mean_test_score', 'std_test_score', 'params']]
3 a['C'] = [d.get('C') for d in a['params']]
4 b=a.sort_values(['C'])
5 CV_Error=1-b['mean_test_score']
6 C =b['C']
7
8
9 plt.plot(C,CV_Error)
10 plt.xlabel('C')
11 plt.ylabel('Cross-Validated Error')
```

```
Text(0,0.5,'Cross-Validated Error')
```



```
1
2 #{'class_weight': 'balanced', 'C': 50}
3 LR_optimal=LogisticRegression(penalty='l2',C=50,class_weight='balanced')
4
5 # fitting the model
6 LR_optimal.fit(sent_vectors_train_avg2v, y_train)
7
8 # predict the response
9 pred_avg_w2v = LR_optimal.predict(sent_vectors_test_avg2v)
10
11 # evaluate f1_score
12 f1_score = f1_score(y_test, pred_avg_w2v)
13
14 # Train & Test Error
15 print("The overall f1_score for the Train Data is : ", metrics.f1_score(y_train,LR_optimal.predict(sent_
16 print("The overall f1_score for the Test Data is : ", metrics.f1_score(y_test,pred_avg_w2v))
```


The overall f1\_score for the Train Data is : 0.5152863591253961

The overall f1\_score for the Test Data is : 0.535966149506347


```
1 #Confusion matrix
2 C = confusion_matrix(y_test, pred_avg_w2v)
3 A =(((C.T)/(C.sum(axis=1))).T)
4 B =(C/C.sum(axis=0))
5 labels = [0,1]
```

```
1  print("-"*20, "Confusion matrix", "-"*20)
2  plt.figure(figsize=(12,7))
3  sns.heatmap(C, annot=True, cmap="YlGnBu", fmt="g", xticklabels=labels, yticklabels=labels)
4  plt.xlabel('Predicted Class')
5  plt.ylabel('Original Class')
6  plt.show()
7
8  print("-"*20, "Precision matrix (Column Sum=1)", "-"*20)
9  plt.figure(figsize=(12,7))
10 sns.heatmap(B, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels, yticklabels=labels)
11  plt.xlabel('Predicted Class')
12  plt.ylabel('Original Class')
13  plt.show()
14
15  # representing B in heatmap format
16  print("-"*20, "Recall matrix (Row sum=1)", "-"*20)
17  plt.figure(figsize=(12,7))
18  sns.heatmap(A, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels, yticklabels=labels)
19  plt.xlabel('Predicted Class')
20  plt.ylabel('Original Class')
21  plt.show()
```


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



### TF-Idf W2Vec

```
1 tf_idf_vect = TfidfVectorizer()
2 vocabulary = tf_idf_vect.fit(train_data['CleanedText'])
3 final_tf_idf= tf_idf_vect.transform(train_data['CleanedText'])
4
5 # we are converting a dictionary with word as a key, and the idf as a value
6 dictionary = dict(zip(vocabulary.get_feature_names(), list(tf_idf_vect.idf_)))
```








```
1 grid.fit(tfidf_w2v_sent_vectors_train, y_train)
2
3 # examine the best model
4 print(grid.best_score_)
5 print(grid.best_params_)
```

0.47618209651423954

{'class\_weight': 'balanced', 'C': 0.5}

```
1 #Plotting C v/s CV_error
2 a=pd.DataFrame(grid.cv_results_)[['mean_test_score', 'std_test_score', 'params']]
3 a['C'] = [d.get('C') for d in a['params']]
4 b=a.sort_values(['C'])
5 CV_Error=1-b['mean_test_score']
6 C =b['C']
7
8
9 plt.plot(C,CV_Error)
10 plt.xlabel('C')
11 plt.ylabel('Cross-Validated Error')
```

```
Text(0,0.5,'Cross-Validated Error')
```



```
1  #{'class_weight': 'balanced', 'C': 0.5}
2  LR_optimal=LogisticRegression(penalty='l2',C=0.5,class_weight='balanced')
3
4  # fitting the model
5  LR_optimal.fit(tfidf_w2v_sent_vectors_train, y_train)
6
7  # predict the response
8  pred_tfidf_w2v_sent_vectors_test = LR_optimal.predict(tfidf_w2v_sent_vectors_test)
9
10 # evaluate f1_score
11 f1_score = f1_score(y_test, pred_tfidf_w2v_sent_vectors_test)
12
13 # Train & Test Error
14 print("The overall f1_score for the Train Data is : ", metrics.f1_score(y_train,LR_optimal.predict(tfidf_w2v_sent_vectors_train)))
15 print("The overall f1_score for the Test Data is : ", metrics.f1_score(y_test,pred_tfidf_w2v_sent_vectors_test))
```


The overall f1\_score for the Train Data is : 0.47774873135475937

The overall f1\_score for the Test Data is : 0.5006821282401092


```
1  #Confusion matrix
2  C = confusion_matrix(y_test, pred_tfidf_w2v_sent_vectors_test)
3  A = (((C.T)/(C.sum(axis=1))).T)
4  B = (C/C.sum(axis=0))
5  labels = [0,1]
```

```
1  print("-"*20, "Confusion matrix", "-"*20)
2  plt.figure(figsize=(12,7))
3  sns.heatmap(C, annot=True, cmap="YlGnBu", fmt="g", xticklabels=labels, yticklabels=labels)
4  plt.xlabel('Predicted Class')
5  plt.ylabel('Original Class')
6  plt.show()
7
8  print("-"*20, "Precision matrix (Column Sum=1)", "-"*20)
9  plt.figure(figsize=(12,7))
10 sns.heatmap(B, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels, yticklabels=labels)
11 plt.xlabel('Predicted Class')
12 plt.ylabel('Original Class')
13 plt.show()
14
15     # representing B in heatmap format
16 print("-"*20, "Recall matrix (Row sum=1)", "-"*20)
17 plt.figure(figsize=(12,7))
18 sns.heatmap(A, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels, yticklabels=labels)
19 plt.xlabel('Predicted Class')
20 plt.ylabel('Original Class')
21 plt.show()
```


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



Reporting f1\_score for above featurization with L2 regularizer

```
1 from prettytable import PrettyTable
```



```
1 x=PrettyTable()
2 x.field_names = ["Model", "Bow", "Tfidf", "Avg-W2V", "Tfidf-W2V"]
3 x.add_row(["C", 0.01, 1e-07, 50, 0.5])
4 x.add_row(["Train f1_score", 0.80, 0.90, 0.51, 0.47])
5 x.add_row(["Test f1_score", 0.67, 0.53, 0.53, 0.50])
6
7 print(x)
```

```
+-----+-----+-----+-----+
|   Model   | Bow | Tfidf | Avg-W2V | Tfidf-W2V |
+-----+-----+-----+-----+
|      C      | 0.01 | 1e-07 | 50      | 0.5       |
| Train f1_score | 0.8  | 0.9   | 0.51    | 0.47     |
| Test f1_score  | 0.67 | 0.53  | 0.53    | 0.5       |
+-----+-----+-----+-----+
```